

Table of Contents

1. Syntax for outer joins (Left, Right, Full).....	1
2. Queries using altgeld_mart tables	3
2.1. Customers and orders	3
2.2. Products and orders	5

These are the tables we are using. Note that we have employees with no projects and a department with no employees and employees with no department.

z_em_dept		z_em_emp			z_em_empproj	
d_id	d_name	e_id	e_name	d_id	p_id	e_id
100	Manufacturing	1	Jones	150	ORDB-10	3
150	Accounting	2	Martin	150	ORDB-10	5
200	Marketing	3	Gates	250	Q4-SALES	2
250	Research	4	Anders	100	Q4-SALES	4
		5	Bossy		ORDB-10	2
		6	Perkins		Q4-SALES	5

1. Syntax for outer joins (Left, Right, Full)

Outer joins can use the syntax Left Join or Right Join. A join written as

```
From tblA LEFT JOIN tblB
```

will include all rows from table tblA and any matching rows from tblB. The table to the left of the phrase Left Join will have all of its rows returned.

A join written as

```
From tblA RIGHT JOIN tblB
```

will include all rows from table tblB and any matching rows from tblA. The table to the right of the phrase Right Join will have all of its rows returned.

The outer joins are not symmetric.

The word OUTER is optional; you can use Left Outer Join or LeftJoin.

You will still need to identify the joining columns.

Demo 01: All departments; employees of those departments if they exist.

```
Select z_em_dept.d_id, d_name, e_id, e_name
From z_em_dept
LEFT JOIN z_em_emp on z_em_dept.d_id = z_em_emp.d_id
;
```

d_id	d_name	e_id	e_name
100	Manufacturing	4	Anders
150	Accounting	1	Jones
150	Accounting	2	Martin
200	Marketing	NULL	NULL
250	Research	3	Gates

Demo 02: All employees; assigned departments if they exist. Outer joins are not commutative

```

Select d_id, d_name, e_id, e_name
From z_em_emp
LEFT JOIN z_em_dept using(d_id)
;

```

d_id	d_name	e_id	e_name
150	Accounting	1	Jones
150	Accounting	2	Martin
250	Research	3	Gates
100	Manufacturing	4	Anders
NULL	NULL	5	Bossy
NULL	NULL	6	Perkins

Demo 03: All employees; assigned departments if they exist.

```

Select d_id, d_name, e_id, e_name
From z_em_dept
RIGHT JOIN z_em_emp using(d_id)
;

```

d_id	d_name	e_id	e_name
150	Accounting	1	Jones
150	Accounting	2	Martin
250	Research	3	Gates
100	Manufacturing	4	Anders
NULL	NULL	5	Bossy
NULL	NULL	6	Perkins

MySQL does not yet support the Full Outer join syntax that you might know from another dbms.

Demo 04: Three table outer join. This is all of the departments and their employees if there are any in the department and the projects if the employees have a project.

```

Select d_id, d_name, e_id, e_name, p_id
From z_em_dept
LEFT JOIN z_em_emp using(d_id)
LEFT JOIN z_em_empproj using(e_id);

```

d_id	d_name	e_id	e_name	p_id
100	Manufacturing	4	Anders	Q4-SALES
150	Accounting	1	Jones	NULL
150	Accounting	2	Martin	Q4-SALES
150	Accounting	2	Martin	ORDB-10
200	Marketing	NULL	NULL	NULL
250	Research	3	Gates	ORDB-10

Demo 05: Three table outer join. This is all of the employees and their departments if they have one and their projects if they have one

```
Select e_id, e_name, d_id, d_name, p_id
From z_em_emp
LEFT JOIN z_em_dept using(d_id)
LEFT JOIN z_em_empproj using(e_id)
Order by z_em_emp.e_id;
```

e_id	e_name	d_id	d_name	p_id
1	Jones	150	Accounting	NULL
2	Martin	150	Accounting	Q4-SALES
2	Martin	150	Accounting	ORDB-10
3	Gates	250	Research	ORDB-10
4	Anders	100	Manufacturing	Q4-SALES
5	Bossy	NULL	NULL	ORDB-10
5	Bossy	NULL	NULL	Q4-SALES
6	Perkins	NULL	NULL	NULL

Demo 06: Suppose we want to see all employees and their departments if they have one and the names of their projects if they have one. The following query does not do that. We start with an outer join but then use an inner join which eliminates employees with no projects.

```
Select z_em_emp.e_id, e_name, z_em_dept.d_id, d_name, p_id
From z_em_emp
LEFT JOIN z_em_dept on z_em_dept.d_id = z_em_emp.d_id
JOIN      z_em_empproj on z_em_emp.e_id = z_em_empproj.e_id
Order by  z_em_emp.e_id;
```

e_id	e_name	d_id	d_name	p_id
2	Martin	150	Accounting	ORDB-10
2	Martin	150	Accounting	Q4-SALES
3	Gates	250	Research	ORDB-10
4	Anders	100	Manufacturing	Q4-SALES
5	Bossy	NULL	NULL	ORDB-10
5	Bossy	NULL	NULL	Q4-SALES

2. Queries using altgeld_mart tables

2.1. Customers and orders

The cust_id filter is simply to reduce the volume of output.

Demo 07: Customers with orders. This uses an inner join.

```
Select cust_id
, cust_name_last
, ord_id
From a_oe.customers
JOIN a_oe.order_headers using(cust_id)
Where cust_id between 404900 and 409030
Order by cust_id, ord_id
;
```

```

+-----+-----+-----+
| cust_id | cust_name_last | ord_id |
+-----+-----+-----+
| 404900 | Williams       | 520    |
| 404950 | Morris         | 110    |
| 404950 | Morris         | 408    |
| 404950 | Morris         | 510    |
| 404950 | Morris         | 535    |
| 404950 | Morris         | 540    |
| 405000 | Day            | 116    |
| 408770 | Clay           | 405    |
| 409030 | Mazur          | 128    |
| 409030 | Mazur          | 130    |
| 409030 | Mazur          | 324    |
+-----+-----+-----+
11 rows in set (0.03 sec)

```

Demo 08: Customers with and without orders. This uses an outer join; Customers Left Join Order Headers. That means we get customers with orders and if the customer has several orders, that customer gets multiple lines in the result set. We also get rows for the two customers in this cust_id range who have no orders and the column for their order id value is null- these customers each get one row.

```

Select cust_id
, cust_name_last
, ord_id
From a_oe.customers
LEFT JOIN a_oe.order_headers using(cust_id)
Where cust_id between 404900 and 409030
Order by cust_id, ord_id;
+-----+-----+-----+
| cust_id | cust_name_last | ord_id |
+-----+-----+-----+
| 404900 | Williams       | 520    |
| 404950 | Morris         | 110    |
| 404950 | Morris         | 408    |
| 404950 | Morris         | 510    |
| 404950 | Morris         | 535    |
| 404950 | Morris         | 540    |
| 405000 | Day            | 116    |
| 408770 | Clay           | 405    |
| 409010 | Morris         | NULL   |
| 409020 | Max            | NULL   |
| 409030 | Mazur          | 128    |
| 409030 | Mazur          | 130    |
| 409030 | Mazur          | 324    |
+-----+-----+-----+
13 rows in set (0.03 sec)

```

Demo 09: Now consider this join. I change the join to a right join. The result set is the same as the inner join used previously. Why?

```

Select cust_id
, cust_name_last
, ord_id
From a_oe.customers
RIGHT JOIN a_oe.order_headers using(cust_id)

```

```
Where cust_id between 404900 and 409030
Order by cust_id, ord_id;
```

```
+-----+-----+-----+
| cust_id | cust_name_last | ord_id |
+-----+-----+-----+
| 404900 | Williams       | 520    |
| 404950 | Morris         | 110    |
| 404950 | Morris         | 408    |
| 404950 | Morris         | 510    |
| 404950 | Morris         | 535    |
| 404950 | Morris         | 540    |
| 405000 | Day            | 116    |
| 408770 | Clay           | 405    |
| 409030 | Mazur          | 128    |
| 409030 | Mazur          | 130    |
| 409030 | Mazur          | 324    |
+-----+-----+-----+
11 rows in set (0.03 sec)
```

In our database we have a foreign key in the order headers table that refers back to the customer table and to the `cust_id` in the customer table.

```
create table a_oe.order_headers(
    ord_id          int unsigned not null
  , ord_date        datetime      not null
  , cust_id         int unsigned not null
  . . .
  , constraint ord_cust_fk foreign key(cust_id) references a_oe.customers(cust_id)
  . . . )
```

I also set the `cust_id` in the order headers table as Not null. This means that every row in the order headers table must have a value for the `cust_id` (it is Not null) and that `cust_id` in the order header must match a `cust_id` in the customers tables (foreign key reference).

The outer join in this query is asking for all orders whether or not they match a customer. But our database is set up so that every order header rows is matched with a customer. So it does not make sense to ask to see order headers rows that do not match a customer. In this case you should use an inner join. Using an outer join when it is logically impossible to return unmatched rows is inefficient. Someone reading your query would assume you have made a mistake someplace but they would not know what the mistake is- is the database badly designed and allows the entry of orders that do not belong to a customer (who pays for those orders?), or did you get the join order incorrect?

2.2. Products and orders

These are limited to products in the MUS category to reduce the volume of output

Demo 10: First an inner join- these show products which have been ordered- each product id must match a product id on an order detail row

```
Select PR.prod_id, PR.prod_desc, PR.catg_id, OD.ord_id
From a_prd.products PR
JOIN a_oe.order_details OD on PR.prod_id = OD.prod_id
Where PR.catg_id in ('MUS')
Order by PR.prod_id;
```

```
+-----+-----+-----+-----+
| prod_id | prod_desc                | catg_id | ord_id |
+-----+-----+-----+-----+
| 2014    | Bix Beiderbecke - Tiger Rag | MUS     | 518    |
| 2014    | Bix Beiderbecke - Tiger Rag | MUS     | 525    |
```

```

| 2014 | Bix Beiderbecke - Tiger Rag | MUS | 715 |
| 2412 | David Newman - Davey Blue | MUS | 525 |
| 2746 | Charles Mingus - Blues & Politics | MUS | 525 |
| 2747 | Charles Mingus - Blues & Roots | MUS | 520 |
| 2947 | Ornette Coleman - Sound Grammer | MUS | 525 |
| 2984 | John Coltrane - Lush Life | MUS | 518 |
| 2984 | John Coltrane - Lush Life | MUS | 715 |
+-----+-----+-----+-----+
9 rows in set (0.00 sec)

```

Demo 11: How many products do we have in the MUS category?
 We have 11 products; looking at the previous result set, 6 of these products were sold (One was on two different orders)

```

Select PR.prod_id, PR.prod_desc, PR.catg_id
From a_prd.products PR
Where catg_id in ('MUS')
Order by PR.prod_id;
+-----+-----+-----+-----+
| prod_id | prod_desc | catg_id |
+-----+-----+-----+-----+
| 2014 | Bix Beiderbecke - Tiger Rag | MUS |
| 2234 | Charles Mingus - Pithecanthropus Erectus | MUS |
| 2337 | John Coltrane - Blue Train | MUS |
| 2412 | David Newman - Davey Blue | MUS |
| 2487 | Stanley Turrentine - Don't Mess With Mr. T | MUS |
| 2746 | Charles Mingus - Blues & Politics | MUS |
| 2747 | Charles Mingus - Blues & Roots | MUS |
| 2933 | David Newman - I Remember Brother Ray | MUS |
| 2947 | Ornette Coleman - Sound Grammer | MUS |
| 2984 | John Coltrane - Lush Life | MUS |
| 2987 | Stanley Turrentine - Ballads | MUS |
+-----+-----+-----+-----+
11 rows in set (0.00 sec)

```

Demo 12: We can use an outer join to get both ordered and un-ordered products. Why do we get 12 rows?

```

Select PR.prod_id, prod_desc, catg_id, ord_id
From a_prd.products PR
LEFT JOIN a_oe.order_details OD on PR.prod_id = OD.prod_id
Where catg_id in ('MUS')
Order by PR.prod_id;
+-----+-----+-----+-----+
| prod_id | prod_desc | catg_id | ord_id |
+-----+-----+-----+-----+
| 2014 | Bix Beiderbecke - Tiger Rag | MUS | 518 |
| 2014 | Bix Beiderbecke - Tiger Rag | MUS | 525 |
| 2014 | Bix Beiderbecke - Tiger Rag | MUS | 715 |
| 2234 | Charles Mingus - Pithecanthropus Erectus | MUS | NULL |
| 2337 | John Coltrane - Blue Train | MUS | NULL |
| 2412 | David Newman - Davey Blue | MUS | 525 |
| 2487 | Stanley Turrentine - Don't Mess With Mr. T | MUS | NULL |
| 2746 | Charles Mingus - Blues & Politics | MUS | 525 |
| 2747 | Charles Mingus - Blues & Roots | MUS | 520 |
| 2933 | David Newman - I Remember Brother Ray | MUS | NULL |
| 2947 | Ornette Coleman - Sound Grammer | MUS | 525 |
| 2984 | John Coltrane - Lush Life | MUS | 518 |
| 2984 | John Coltrane - Lush Life | MUS | 715 |
+-----+-----+-----+-----+

```

```

| 2987 | Stanley Turrentine - Ballads | MUS | NULL |
+-----+-----+-----+
14 rows in set (0.00 sec)

```

Demo 13: This query gives us rows for the same products- why are we missing values in the first column which shows the product id? Every product has a product Id!

```

Select OD.prod_id, prod_desc, catg_id, ord_id
From a_prd.products PR
LEFT JOIN a_oe.order_details OD on PR.prod_id = OD.prod_id
Where catg_id in ('MUS')
Order by OD.prod_id;
+-----+-----+-----+
| prod_id | prod_desc | catg_id | ord_id |
+-----+-----+-----+
| NULL | John Coltrane - Blue Train | MUS | NULL |
| NULL | David Newman - I Remember Brother Ray | MUS | NULL |
| NULL | Stanley Turrentine - Ballads | MUS | NULL |
| NULL | Stanley Turrentine - Don't Mess With Mr. T | MUS | NULL |
| NULL | Charles Mingus - Pithecanthropus Erectus | MUS | NULL |
| 2014 | Bix Beiderbecke - Tiger Rag | MUS | 525 |
| 2014 | Bix Beiderbecke - Tiger Rag | MUS | 715 |
| 2014 | Bix Beiderbecke - Tiger Rag | MUS | 518 |
| 2412 | David Newman - Davey Blue | MUS | 525 |
| 2746 | Charles Mingus - Blues & Politics | MUS | 525 |
| 2747 | Charles Mingus - Blues & Roots | MUS | 520 |
| 2947 | Ornette Coleman - Sound Grammer | MUS | 525 |
| 2984 | John Coltrane - Lush Life | MUS | 715 |
| 2984 | John Coltrane - Lush Life | MUS | 518 |
+-----+-----+-----+
14 rows in set (0.00 sec)

```

What I did is switch the column alias for the first column and for the sort key to use the order details table. If I am looking for the product id in the order details table, the products which are not ordered do not have a value for that column and display as nulls.