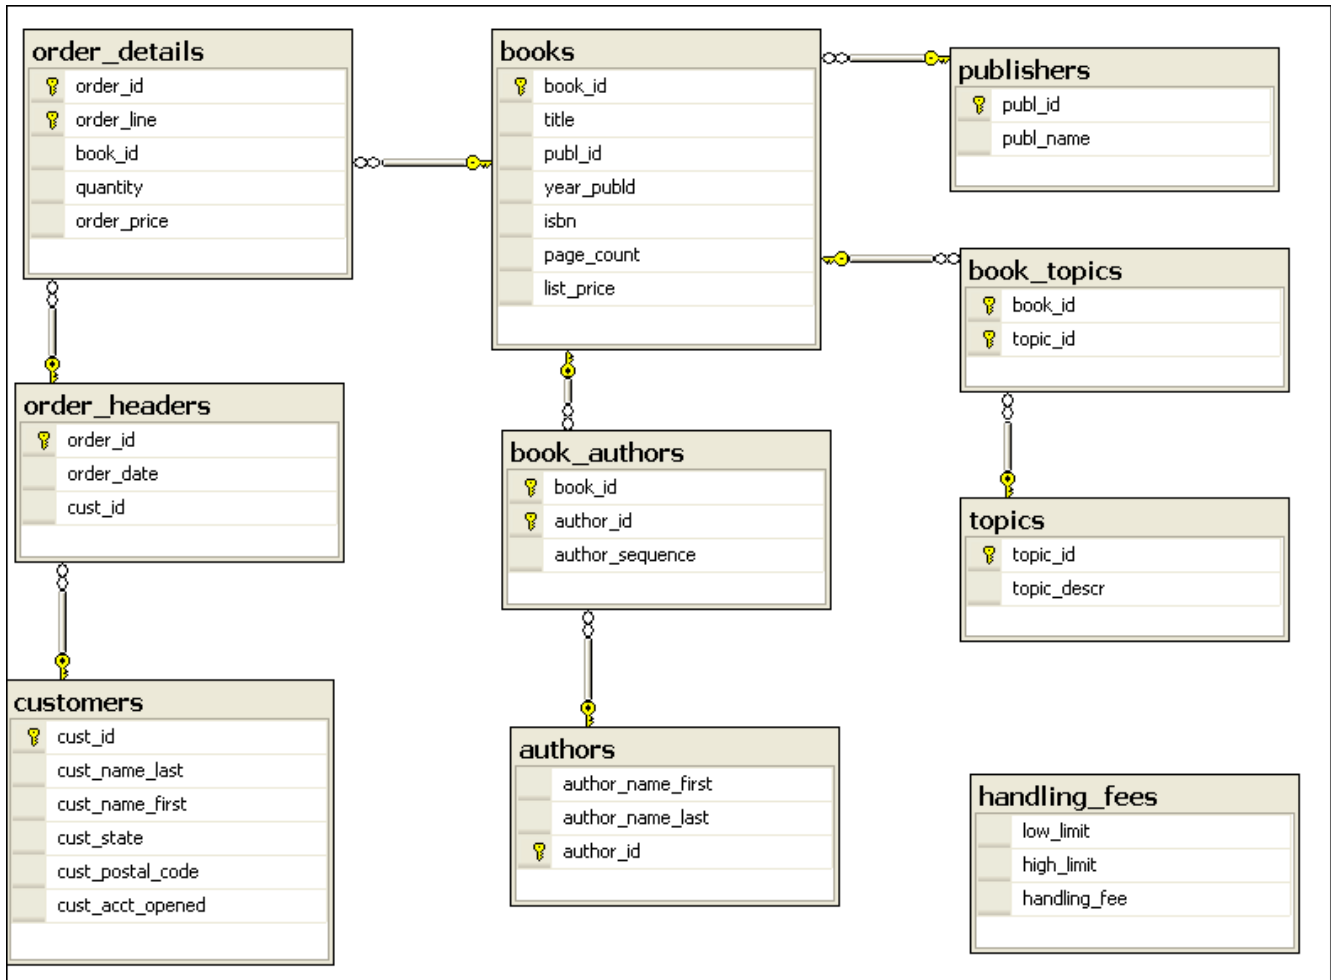


The books database is used to store data for a small bookstore.

These are the tables for the books database. Try reading through the create table statements to see how the tables are related to each other. You could experiment with some select statements and try a few joins. Experimenting is a good way to get a feel for these tables.



We have a books table containing basic data about the books we carry. The tables each have a primary key, except for the handling\_fees table. You should read the create table statements to find how these rules are described in SQL.

### Books table

We have several attributes for books. The primary key is the book\_id. The title is required. The publisher id is nullable. This means we could carry books that are not published by an established publisher. The year published is a required field and we do not carry any book published before 1850. The ISBN is nullable- some books might not have an ISBN:

We have a list\_price; a book can be free; this is a nullable field which means we could enter a book for which we do not know the list price and we could enter a book with a list\_price of 0. The page count has to be greater than or equal to 0 and is also nullable.

A book is associated with only one publisher.

**Publishers table:**

We are storing only the publisher name. This is used as a lookup table so that we can store the shorter publ\_id values in the books table. The publisher name is required. We could also store more information about publishers in the table- such as mailing address, phone numbers for our contact person at the publisher.

**Categories of Books:**

Books can be classified under multiple categories - so a book that compared Oracle SQL to T-SQL could be classified under the categories SQL, ORA and SSRV.

**Topics table:**

In this table we have the short topic\_id field and the longer descriptive field. This is used as a lookup table so that we can store the shorter topic\_id in the books table. The descriptive field is unique and not null; we do not want to have two different topics with the same descriptive name

**Book Topics table:**

Since a book can have multiple topics and a topic can be used for multiple books, we use a junction table bookTopics to handle this. This table has only two columns- one to identify the book and the other to identify the topic. These columns are fk to the books and topics table. Together they make up the pk. This means that a book which has one topic would have one row in the bookTopics table while a book which has three topics would have three rows in the bookTopics table. It is possible that a book has no rows in the bookTopics table if it does not fit in one of our predefined topic area. And we can have set up topic areas and not have any books classified as that topic.

This table is said to be "all-keys" since all of its columns are FK to other tables and make up the PK.

A book could have multiple authors and an author could have written multiple books.

**Authors table:**

We are storing the authors first and last name.

**Book authors table**

This is similar in nature to the bookTopics table since it joins the books and authors tables which have a M:N relationship. It also includes a column, author sequence, which indicates the order of the listing of the authors of the book. Authors care about the order in which their name is listed.

These are some simplifying assumptions we are making about books. We are not keeping information about the various editions of a book or that the same book might have been published in different formats over the years (hardcopy , ebook, books on tape) or have been taken over by different publishers. We assume each of these variations would get its own book id and we are not tracking that these rows refer to essentially the same book. And these are not rare books where each individual copy of a book would have its own row.

We are storing only a few fields for customers (mostly to keep the data table smaller). We have the customer's name, state, and postal code and the date they opened an account with us.

Customers can place orders for books and each order belongs to only one customer. An order can be for one or more books and a customer can order multiple copies of a book. This database is not concerned with the cash type sales that might occur at a bricks and mortar store.

**customers, order headers, order details tables:**

These are standard table structures. Read the sql for these. The attributes in the order tables are not null. If we delete a row in the order table that also deletes the related rows in the order details table. The OrderDetails table has a compound primary key.

**handling\_fees**

This is the one table that does not have a primary key. It is a range lookup table. Suppose we are shipping an order with 10 books. We can find which range in the handling fees lookup table is matched and can determine the fee. One issue with this type of table is the need for a value for the highest volume. We have an assumption here that we will never ship more than 999 books on an order. That assumption could cause us problems.

We could set the low\_limit column as a pk but that is less meaningful than the pk columns for the other tables.